

CSU22011: ALGORITHMS AND DATA STRUCTURES I

Lecture 2: Analysis of Algorithms

Vasileios Koutavas



School of Computer Science and Statistics
Trinity College Dublin

Programs = Algorithms + Data structures.

→ Must be **correct**

→ Must be **efficient**

Programs = Algorithms + Data structures.

→ Must be **correct**

→ Must be **efficient**

→ What is efficiency?

→ **Analysis of Algorithms**: How to evaluate efficiency?

Programs = Algorithms + Data structures.

→ Must be **correct**

→ Must be **efficient**

→ What is efficiency?

→ **Analysis of Algorithms**: How to evaluate efficiency?

→ Established measure of efficiency (from now on **performance**):
how the program **scales** to larger inputs. Measure the effect of
doubling the input size to

→ the **running time** of the algorithm

→ the **memory footprint** of the algorithm

Programs = Algorithms + Data structures.

→ Must be **correct**

→ Must be **efficient**

→ What is efficiency?

→ **Analysis of Algorithms**: How to evaluate efficiency?

→ Established measure of efficiency (from now on **performance**):
how the program **scales** to larger inputs. Measure the effect of
doubling the input size to

→ the **running time** of the algorithm

→ the **memory footprint** of the algorithm

→ We want to be able to **analyse algorithms** w.r.t. their performance.

WHY ANALYSE ALGORITHMS?

- **Good programmer:** to predict the performance of our programs.
- **Good client:** to choose between alternative algorithms/implementations.
- **Good manager:** to provide guarantees to clients / avoid client complaints.
- **Good scientist:** to understand the nature of computing.

EXAMPLE: LINEAR SEARCH

Linear Search: simple search in an array:

```
boolean linearSearch1(int[] ar, int s) {  
    boolean found = false;  
    for (int i = 0; i < ar.length; i++) {  
        if (ar[i] == s) found = true;  
    }  
    return found;  
}
```

How can we evaluate how well the running time of this algorithm scales to larger inputs?

¹We will see a number of approaches how to evaluate running times

EXAMPLE: LINEAR SEARCH

Linear Search: simple search in an array:

```
boolean linearSearch1(int[] ar, int s) {  
    boolean found = false;  
    for (int i = 0; i < ar.length; i++) {  
        if (ar[i] == s) found = true;  
    }  
    return found;  
}
```

How can we evaluate how well the running time of this algorithm scales to larger inputs?

- Evaluate¹ its running time for various sizes of array **ar**.
Then calculate the **rate of growth** of the running time with respect to the input size (stay tuned).

¹We will see a number of approaches how to evaluate running times

EXAMPLE: LINEAR SEARCH 2

Improved Linear Search: return as soon as possible.

```
boolean linearSearch2(int[] ar, int s) {  
    for (int i = 0; i < ar.length; i++) {  
        if (ar[i] == s) return true;  
    }  
    return false;  
}
```

Consider inputs arrays of size 100. Is the running time the same?

EXAMPLE: LINEAR SEARCH 2

Improved Linear Search: return as soon as possible.

```
boolean linearSearch2(int[] ar, int s) {  
    for (int i = 0; i < ar.length; i++) {  
        if (ar[i] == s) return true;  
    }  
    return false;  
}
```

Consider inputs arrays of size 100. Is the running time the same?

NO! Even for the same input size, some inputs will require more time than others.

- best case inputs; e.g: $\text{ar}=[1,2,\dots,100]$, $s = 1$
- **worst case** inputs; e.g: $\text{ar}=[1,2,\dots,100]$, $s = 101$
- everything in between (sometimes we talk about average case)

EXAMPLE: LINEAR SEARCH 2

Improved Linear Search: return as soon as possible.

```
boolean linearSearch2(int[] ar, int s) {  
    for (int i = 0; i < ar.length; i++) {  
        if (ar[i] == s) return true;  
    }  
    return false;  
}
```

Consider inputs arrays of size 100. Is the running time the same?


NO! Even for the same input size, some inputs will require more time than others.

- best case inputs; e.g: `ar=[1,2,...,100]`, `s = 1`
- **worst case** inputs; e.g: `ar=[1,2,...,100]`, `s = 101`
- everything in between (sometimes we talk about average case)

In this module (and usually in programming practice) we focus on the **worst-case inputs**.

In the course of the next lectures we will see the following methods for evaluating how an algorithm's running time scales:


1. **Scientific method:** measure running times through experiments
2. **Mathematical methods:** consider a convenient model of computation and:
 - sum up the number of program steps for **worst-case** inputs of size N (N is a variable).
 - calculate the number of program steps for **worst-case** inputs of size N , when N nears infinity.This is the **asymptotic** (worst-case) running time — notation big-O, Ω , Θ .



- Parts from S&W 1.4
- Estimate the performance of algorithms by
 - **Experiments** & Observations
 - **Precise** Mathematical Calculations

Running time

“ As soon as an Analytic Engine exists, it will necessarily guide the future course of the science. Whenever any result is sought by its aid, the question will arise—By what course of calculation can these results be arrived at by the machine in the shortest time? ” — *Charles Babbage (1864)*



how many times do you
have to turn the crank?

Analytic Engine


Some algorithmic successes

Discrete Fourier transform.

- Break down waveform of N samples into periodic components.
- Applications: DVD, JPEG, MRI, astrophysics,
- Brute force: N^2 steps.
- FFT algorithm: $N \log N$ steps, **enables new technology.**



Friedrich Gauss
1805





Some algorithmic successes

N-body simulation.

- Simulate gravitational interactions among N bodies.
- Brute force: N^2 steps.
- Barnes-Hut algorithm: $N \log N$ steps, enables new research.



Andrew Appel
PU '81



The challenge

Q. Will my program be able to solve a large practical input?

Why is my program so slow ?

Why does it run out of memory ?



Insight. [Knuth 1970s] Use **scientific method** to understand performance.

Scientific method applied to analysis of algorithms

A framework for predicting performance and comparing algorithms.


Scientific method.

- **Observe** some feature of the natural world.
- **Hypothesize** a model that is consistent with the observations.
- **Predict** events using the hypothesis.
- **Verify** the predictions by making further observations.
- **Validate** by repeating until the hypothesis and observations agree.

Principles.

- Experiments must be **reproducible**.
- Hypotheses must be **falsifiable**.

Feature of the natural world. Computer itself.



SCIENTIFIC APPROACH:

EVALUATING PERFORMANCE BY
EXPERIMENTS


Example: 3-SUM

3-SUM. Given N distinct integers, how many triples sum to exactly zero?

```
% more 8ints.txt
8
30 -40 -20 -10 40 0 10 5

% java ThreeSum 8ints.txt
4
```

	a[i]	a[j]	a[k]	sum
1	30	-40	10	0
2	30	-20	-10	0
3	-40	40	0	0
4	-10	0	10	0



Context. Deeply related to problems in computational geometry.

3-SUM: brute-force algorithm

```
public class ThreeSum
{
    public static int count(int[] a)
    {
        int N = a.length;
        int count = 0;
        for (int i = 0; i < N; i++)
            for (int j = i+1; j < N; j++)
                for (int k = j+1; k < N; k++)
                    if (a[i] + a[j] + a[k] == 0)
                        count++;
        return count;
    }

    public static void main(String[] args)
    {
        In in = new In(args[0]);
        int[] a = in.readAllInts();
        StdOut.println(count(a));
    }
}
```

← check each triple
← for simplicity, ignore integer overflow

QUESTION

Suppose we care only about 100-element arrays.

Q. What is a **worst-case input** for ThreeSum?

```
public class ThreeSum
{
    public static int count(int[] a)
    {
        int N = a.length;
        int count = 0;
        for (int i = 0; i < N; i++)
            for (int j = i+1; j < N; j++)
                for (int k = j+1; k < N; k++)
                    if (a[i] + a[j] + a[k] == 0)
                        count++;
        return count;
    }
}
```

QUESTION

Suppose we care only about 100-element arrays.

Q. What is a **worst-case input** for **ThreeSum**?

A. They are all worst case inputs. For-loops always run to the end.

```
public class ThreeSum
{
    public static int count(int[] a)
    {
        int N = a.length;
        int count = 0;
        for (int i = 0; i < N; i++)
            for (int j = i+1; j < N; j++)
                for (int k = j+1; k < N; k++)
                    if (a[i] + a[j] + a[k] == 0)
                        count++;
        return count;
    }
}
```

A. Manual.



tick tick tick

70

tick tick tick tick tick tick tick
tick tick tick tick tick tick tick
tick tick tick tick tick tick tick

528

A pie chart divided into three equal sectors. Two sectors are colored red, and one sector is white. This represents the fraction $\frac{2}{3}$ of the circle being red.

[illegible]

Measuring the running time

Q. How to time a program?

A. Automatic.

```
public class Stopwatch    (part of stdlib.jar)
```


```
    Stopwatch()            create a new stopwatch
```

```
    double elapsedTime()   time since creation (in seconds)
```

```
public static void main(String[] args)
{
    In in = new In(args[0]);
    int[] a = in.readAllInts();
    Stopwatch stopwatch = new Stopwatch();
    StdOut.println(ThreeSum.count(a));
    double time = stopwatch.elapsedTime();
    StdOut.println("elapsed time " + time);
}
```

Empirical analysis

Run the program for various input sizes and measure running time.




Empirical analysis

Run the program for various input sizes and measure running time.

N	time (seconds) †
250	0.0
500	0.0
1,000	0.1
2,000	0.8
4,000	6.4
8,000	51.1
16,000	?


Data analysis

Standard plot. Plot running time $T(N)$ vs. input size N .



Data analysis

Log-log plot. Plot running time $T(N)$ vs. input size N using **log-log scale**.



$$\lg(T(N)) = b \lg N + c$$

$$b = 2.999$$

$$c = -33.2103$$

$$T(N) = a N^b, \text{ where } a = 2^c$$

Regression. Fit straight line through data points: $a N^b$.

Hypothesis. The running time is about $1.006 \times 10^{-10} \times N^{2.999}$ seconds.

power law

slope

Try out the scientific analysis through experiments:

<https://docs.google.com/spreadsheets/d/1WnihyK6g1pYdcT2ndZOqNNRkTitXkWKnOrTgCnM-bw8/edit?usp=sharing>

Prediction and validation

Hypothesis. The running time is about $1.006 \times 10^{-10} \times N^{2.999}$ seconds.



"order of growth" of running time is about N^3 [stay tuned]

Predictions.

- 51.0 seconds for $N = 8,000$.
- 408.1 seconds for $N = 16,000$.

Observations.

N	time (seconds) †
8,000	51.1
8,000	51.0
8,000	51.1
16,000	410.8

validates hypothesis!

Doubling hypothesis

Doubling hypothesis. Quick way to estimate b in a power-law relationship.

Run program, **doubling** the size of the input.

N	time (seconds) †	ratio	lg ratio
250	0.0		–
500	0.0	4.8	2.3
1,000	0.1	6.9	2.8
2,000	0.8	7.7	2.9
4,000	6.4	8.0	3.0
8,000	51.1	8.0	3.0

$$\begin{aligned}\frac{T(2N)}{T(N)} &= \frac{a(2N)^b}{aN^b} \\ &= 2^b\end{aligned}$$

$$\lg(6.4 / 0.8) = 3.0$$

seems to converge to a constant $b \approx 3$

Hypothesis. Running time is about $a N^b$ with $b = \lg \text{ratio}$.

Caveat. Cannot identify logarithmic factors with doubling hypothesis.

Doubling hypothesis

Doubling hypothesis. Quick way to estimate b in a power-law relationship.

Q. How to estimate a (assuming we know b) ?

A. Run the program (for a sufficient large value of N) and solve for a .

N	time (seconds) †
8,000	51.1
8,000	51.0
8,000	51.1

$$51.1 = a \times 8000^3$$

$$\Rightarrow a = 0.998 \times 10^{-10}$$

Hypothesis. Running time is about $0.998 \times 10^{-10} \times N^3$ seconds.



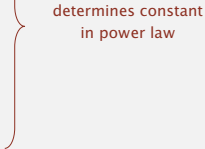
almost identical hypothesis
to one obtained via linear regression

Experimental algorithmics

System independent effects.

- Algorithm.
 - Input data.
- 

System dependent effects.

- Hardware: CPU, memory, cache, ...
 - Software: compiler, interpreter, garbage collector, ...
 - System: operating system, network, other apps, ...
- 

Bad news. Difficult to get precise measurements.

Good news. Much easier and cheaper than other sciences.



e.g., can run huge number of experiments

This was the **scientific approach** to algorithm analysis.

In the **mathematical approach** we do **calculations** instead of experiments.