

# Remember Integration with mutexes?

```
pthread_mutex_t  mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
void *IntegratePart(void *i) {
```

```
    double a,b,area;
```

```
    int rc;
```

```
    a = (int)i * H ;
```

```
    b = a + H;
```

```
    area = trapezoid(a,b);
```

```
    // critical section with mutex
```

```
    rc = pthread_mutex_lock(&mutex);
```

```
    checkResults("pthread_mutex_lock()\n", rc);
```

```
    answer=answer+area;
```

```
    rc = pthread_mutex_unlock(&mutex);
```

```
    checkResults("pthread_mutex_unlock()\n",rc);
```

```
    pthread_exit(NULL);
```

```
}
```

What does Amdahl's law say here?

Assume we have one processor per slice?



Trinity College Dublin

Coláiste na Tríonóide, Baile Átha Cliath

The University of Dublin

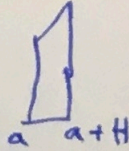
# Some (simplifying) assumptions.

- Assume
  - we can ignore the time spent creating threads
  - assume mutex lock and unlock each require one time-step
  - assume assignment and each arithmetic operation or function call require one time step.:
- Parallelisable code computes a, b, and area (using `trapezoid`, which uses `f`)
  - Cost of `f` is 3
  - Cost of `trapezoid` is 10
- Total cost for a, b and area is 15.
- Serial Code is mutex lock/unlock and answer update: total cost is 4.



# Using a thread to do several trapezoids

$\int_L^H$   
 $n = \text{no of cores}$   
 $w = \text{trapezoids / thread.}$   
 $W = n * w$



$\text{area} = 0;$   
 $\text{launch } n \text{ threads } (t \leftarrow 0 \dots n-1)$  n+1

thread t:

$\text{loc\_area} = 0;$  1  
 $a = L + t * H;$  3  $b = a + H;$  6  
 $\text{for } (i = 0; i \leq w; i++) \{$  1+2w  
      $\text{loc\_area} += \text{trapezoid}(a);$  12 15w  
      $a = b; b = b + H$  2 \*w  
 } 7+17w

lock  
 $\text{area} += \text{loc\_area.}$  xn.  
unlock 4

} Seq: n+1

We get two formulas

Sequential Part:  
 $\text{Seq}(n) = 5n + 1$

Parallel Part:  
 $\text{Par}(w) = 17w + 7$

Time in terms of n, if  
 $W = 1000$ :

$$T(n) = 5n + 8 + 17000/n$$



$$T(n) = 5n + 8 + 17000/n$$

- A bit of calculus shows  $T(n)$  a minimum for  $n = 92$ , when it equals 653 steps
  - $T(1) = 17013$
  - $T(92) = 653$
  - $T(1000) = 5025$
  - $T(3400) = 17013$
- So throwing 3400 processors at this problem is as slow as using one!



# Revisiting Matrix Multiplication

(keeping Amdahl's law in mind)

```
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

What is `p`, the proportion of this program that can be sped-up ?



# We need to consider a specific strategy...

e.g., parallelising just the `k` loop:

```
with k = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,1];  
    c[i,1] = sum;  
    }  
}
```

```
with k = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,2];  
    c[i,2] = sum;  
    }  
}
```

```
with k = 3;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,3];  
    c[i,3] = sum;  
    }  
}
```

What is `p`, the proportion of this program that can be sped-up ?

How do we handle updating `sum`?

Make it a critical resource? Local copies?



# More specific...

parallelising just the `k` loop, with local sums

```
with k = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[1] = 0.0;  
    for j = 1 to 2  
        sum[1] = sum[1] + a[1,j] * b[j,1];  
    c[i,1] = sum[1];  
    }  
}
```

```
with k = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[2] = 0.0;  
    for j = 1 to 2  
        sum[2] = sum[2] + a[1,j] * b[j,2];  
    c[i,2] = sum[2];  
    }  
}
```

```
with k = 3;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[3] = 0.0;  
    for j = 1 to 2  
        sum[3] = sum[3] + a[1,j] * b[j,3];  
    c[i,3] = sum[3];  
    }  
}
```

What is `p`, the proportion of this program that can be sped-up ?

Each `sum[k]` is only used locally!





# Try Six Execution Agents again

```
with k = 1, i = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[1,1] = 0.0;  
    for j = 1 to 2  
        sum[1,1] = sum[1,1] + a[1,j] * b[j,1];  
    c[1,1] = sum[1,1];  
    }  
}
```

...

```
with k = 3, i = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[1,3] = 0.0;  
    for j = 1 to 2  
        sum[1,3] = sum[1,3] + a[1,j] * b[j,3];  
    c[1,3] = sum[1,3];  
    }  
}
```

`sum` is redundant now, and we can simply use `c` itself to compute the result.

```
with k = 1, i = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[2,1] = 0.0;  
    for j = 1 to 2  
        sum[2,1] = sum[2,1] + a[2,j] * b[j,1];  
    c[2,1] = sum[2,1];  
    }  
}
```

...

```
with k = 3, i = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum[2,3] = 0.0;  
    for j = 1 to 2  
        sum[2,3] = sum[2,3] + a[2,j] * b[j,3];  
    c[2,3] = sum[2,3];  
    }  
}
```





# Try Six Execution Agents again

```
with k = 1, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    c[1,1] = 0.0;  
    for j = 1 to 2  
      c[1,1] = c[1,1] + a[1,j] * b[j,1];  
    }  
  }
```

...

```
with k = 3, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    c[1,3] = 0.0;  
    for j = 1 to 2  
      c[1,3] = c[1,3] + a[1,j] * b[j,3];  
    }  
  }
```

```
with k = 1, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    c[2,1] = 0.0;  
    for j = 1 to 2  
      c[2,1] = c[2,1] + a[2,j] * b[j,1];  
    }  
  }
```

...

```
with k = 3, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    c[2,3] = 0.0;  
    for j = 1 to 2  
      c[2,3] = c[2,3] + a[2,j] * b[j,3];  
    }  
  }
```

Here we see we have full output independence between each component of `c`.

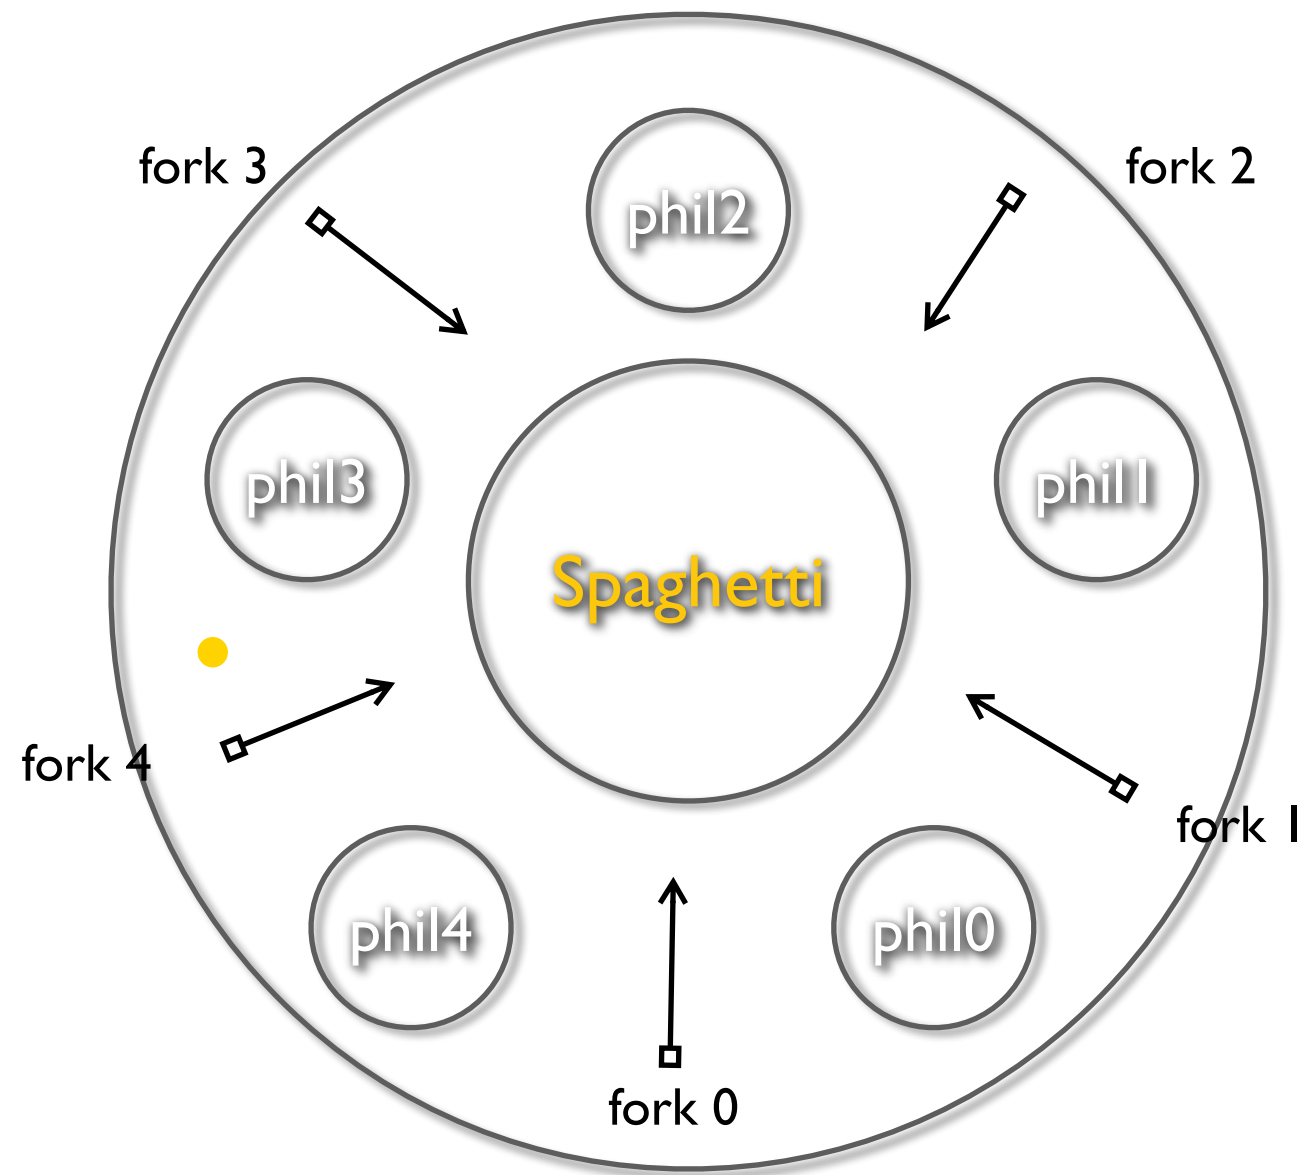


# Matrix Multiply is easy (?)

- Matrix multiply turns out to be highly ("embarrassingly", "massively") parallelisable!
- In principle we should be able to get  $p$  very close to 1.
  - Indeed many supercomputing facilities use massive parallelism to do matrix calculations
- However, for a  $M \times M$  (square) matrix, the equivalent of the "Six Execution Agents" solution requires  $M^2$  processors!
- Also, there are always hidden "non- $p$ " costs in practice
  - Getting fast access by  $N$  cores to main memory gets harder to achieve as  $N$  grows large
  - While each location in matrix  $C$  is only written by one core, each location in both  $A$  and  $B$  are read by many. Again, this can lead to "bus contention" as processors wait to read shared memory.
    - (Clever understanding of cache behaviour can lead to algorithms that minimise such contention)



# Dining Philosophers Problem



Philosophers want to  
repeatedly  
think and then eat.  
Each needs two forks.  
Infinite supply of pasta (ugh!).  
How to ensure they don't starve.

A model problem for thinking about  
problems in Concurrent Programming:  
Deadlock  
Livelock  
Starvation/Fairness  
Data Corruption



# Dining Philosophers

- Features:
  - A philosopher eats only if they have two forks.
  - No two philosophers may hold the same fork simultaneously
- Characteristics of Desired Solution:
  - Freedom from deadlock
  - Freedom from starvation
  - Efficient behaviour generally.



# Modelling the Dining Philosophers

- We imagine the philosophers participate in the following observable events:

Event Shorthand	Description
<i>think.p</i>	Philosopher $p$ is thinking.
<i>eat.p</i>	Philosopher $p$ is eating
<i>pick.p.f</i>	Philosopher $p$ has picked up fork $f$ .
<i>drop.p.f</i>	Philosopher $p$ has dropped fork $f$ .



# What a philosopher does:

- A philosopher wants to: *think ; eat ; think ; eat ; think ; eat ; ....*
- In fact each philosopher needs to do: *think ; pick forks ; eat ; drop forks ; ...*
- We can describe the behaviour of the  $i$ th philosopher as:

$Phil(i) = think.i ; pick.i.i ; pick.i.i+ ; eat.i ; drop.i.i ; drop.i.i+ ; Phil(i)$

- Here  $i+$  is shorthand for  $(i+1) \bmod 5$ .
- What we have are five philosophers running in parallel ( $||$ ):

$Phil(0) || Phil(1) || Phil(2) || Phil(3) || Phil(4)$



# What can (possibly) go wrong ?

- Consider the following (possible, but maybe unlikely) sequence of events, assuming that, just before this point, all philosophers are *think.p*-ing...

*pick.0.0 ; pick.1.1 ; pick.2.2 ; pick.3.3 ; pick 4.4 ;*

- At this point, every philosopher has picked up their left fork.
  - Now each of them wants to pick up its right one.
  - But its right fork is its righthand neighbours left-hand fork!
  - Every philosopher wants to *pick.i.i+* , but can't, because it has already been *pick.i+.i+-ed*!
  - Everyone is stuck and no further progress can be made
- DEADLOCK !





# “Implementing” *pick* and *drop*

- In effect *pick.p.f* attempts to lock a mutex protecting fork *f*.
- So each philosopher is trying to lock two mutexes for two forks before they can *eat.p*.
- The *drop.p.f* simply unlocks the mutex protecting *f*.



# You can't always rely on the scheduler...

- A possible sequence we might observe, starting from when philosophers 1 and 2 are thinking, could be:

*pick.1.1 ; pick.2.2 ; pick.2.3 ; eat.2 ; drop.2.2*

- now, philosopher 1 has picked fork 1 but is waiting for it to be dropped by philosopher 2.
- But philosopher 2 is still running, and so drops the other fork, has a quick think, and then gets quickly back to eating once more:

*pick.1.1 ; pick.2.2 ; pick.2.3 ; eat.2 ; drop.2.2 ; drop.2.3 ; think.2 ; pick.2.2 ; ...*

- Philosopher 1 could get really unlucky and never be scheduled to get the lock for fork 2. It is queuing on the mutex for fork 2, but when philosopher 2 unlocks it, somehow the lock, and control is not handed to philosopher 1.
- STARVATION (and its close friend UN-FAIRNESS)



# Communication

- A concurrent or parallel program consists of two or more separate threads of execution, that run independently except when they have to interact
- To interact, they must *communicate*
- Communication is typically implemented by
  - sharing memory
    - One thread writes data to memory; the other reads it
  - passing messages
    - One thread sends a message; the other gets it



# The Challenge of Concurrency

- Conventional testing and debugging is not generally useful.
  - We assume that the sequential parts of concurrent programs are correctly developed.
- Beyond normal sequential-type bugs, there is a whole range of problems caused by errors in communication and synchronisation. They can not easily be reproduced and corrected.
- So, we are going to use notions, algorithms and formalisms to help us design concurrent programs that are correct by design.



# Sequential Process

- A *sequential process* is the execution of a sequence of *atomic statements*.
  - E.g. Process P could consist of the execution of  $P_1, P_2, P_3, \dots$ .
  - Process Q could be  $Q_1, Q_2, Q_3, \dots$ .
- We think of each sequential process as being a distinct entity that has its own separate program counter (PC).



# Concurrent Execution

- A concurrent system is modelled as a collection of sequential processes, where the atomic statements of the sequential processes can be arbitrarily *interleaved*, but respecting the sequentiality of the atomic statements within their sequential processes.
- E.g. say P is  $P_1, P_2$  and Q is  $Q_1, Q_2$ .



# Scenarios for P and Q.

$p1 \rightarrow q1 \rightarrow p2 \rightarrow q2$
$p1 \rightarrow q1 \rightarrow q2 \rightarrow p2$
$p1 \rightarrow p2 \rightarrow q1 \rightarrow q2$
$q1 \rightarrow p1 \rightarrow q2 \rightarrow p2$
$q1 \rightarrow p1 \rightarrow p2 \rightarrow q2$
$q1 \rightarrow q2 \rightarrow p1 \rightarrow p2$





# Notation

Trivial concurrent program: processes **p** and **q** each do one action that updates global  $n$  with the value of their local variable.

Trivial Concurrent Program (Title)	
integer $n \leftarrow 0$ (Globals)	
p (Process Name)	q
integer $k1 \leftarrow 1$ (Locals)	integer $k2 \leftarrow 2$
p1: $n \leftarrow k1$ (Atomic Statements)	q1: $n \leftarrow k2$

All global and local variables are initialised when declared.



# Simple Sequential Program

## Trivial Sequential Program

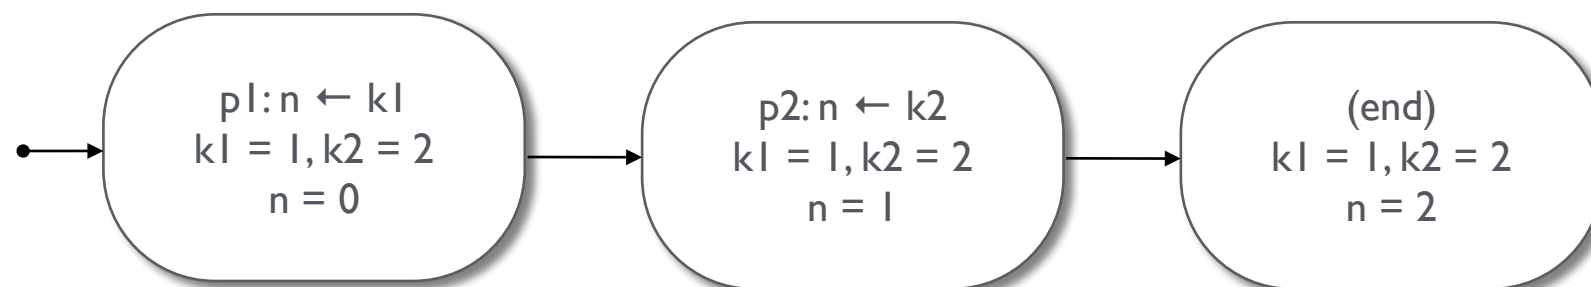
integer  $n \leftarrow 0$

integer  $k1 \leftarrow 1$

integer  $k2 \leftarrow 2$

p1:  $n \leftarrow k1$

p2:  $n \leftarrow k2$


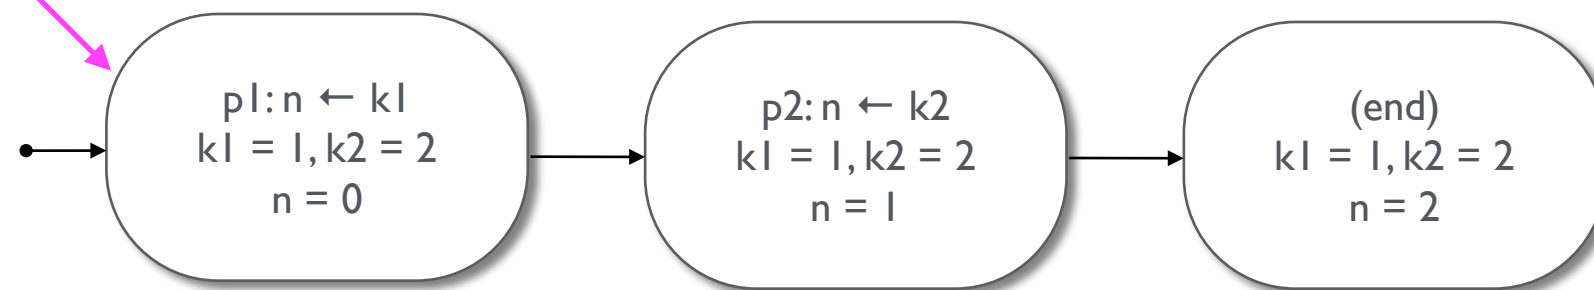


First line gives next atomic action to be executed



# Transition Diagrams

State



Transition



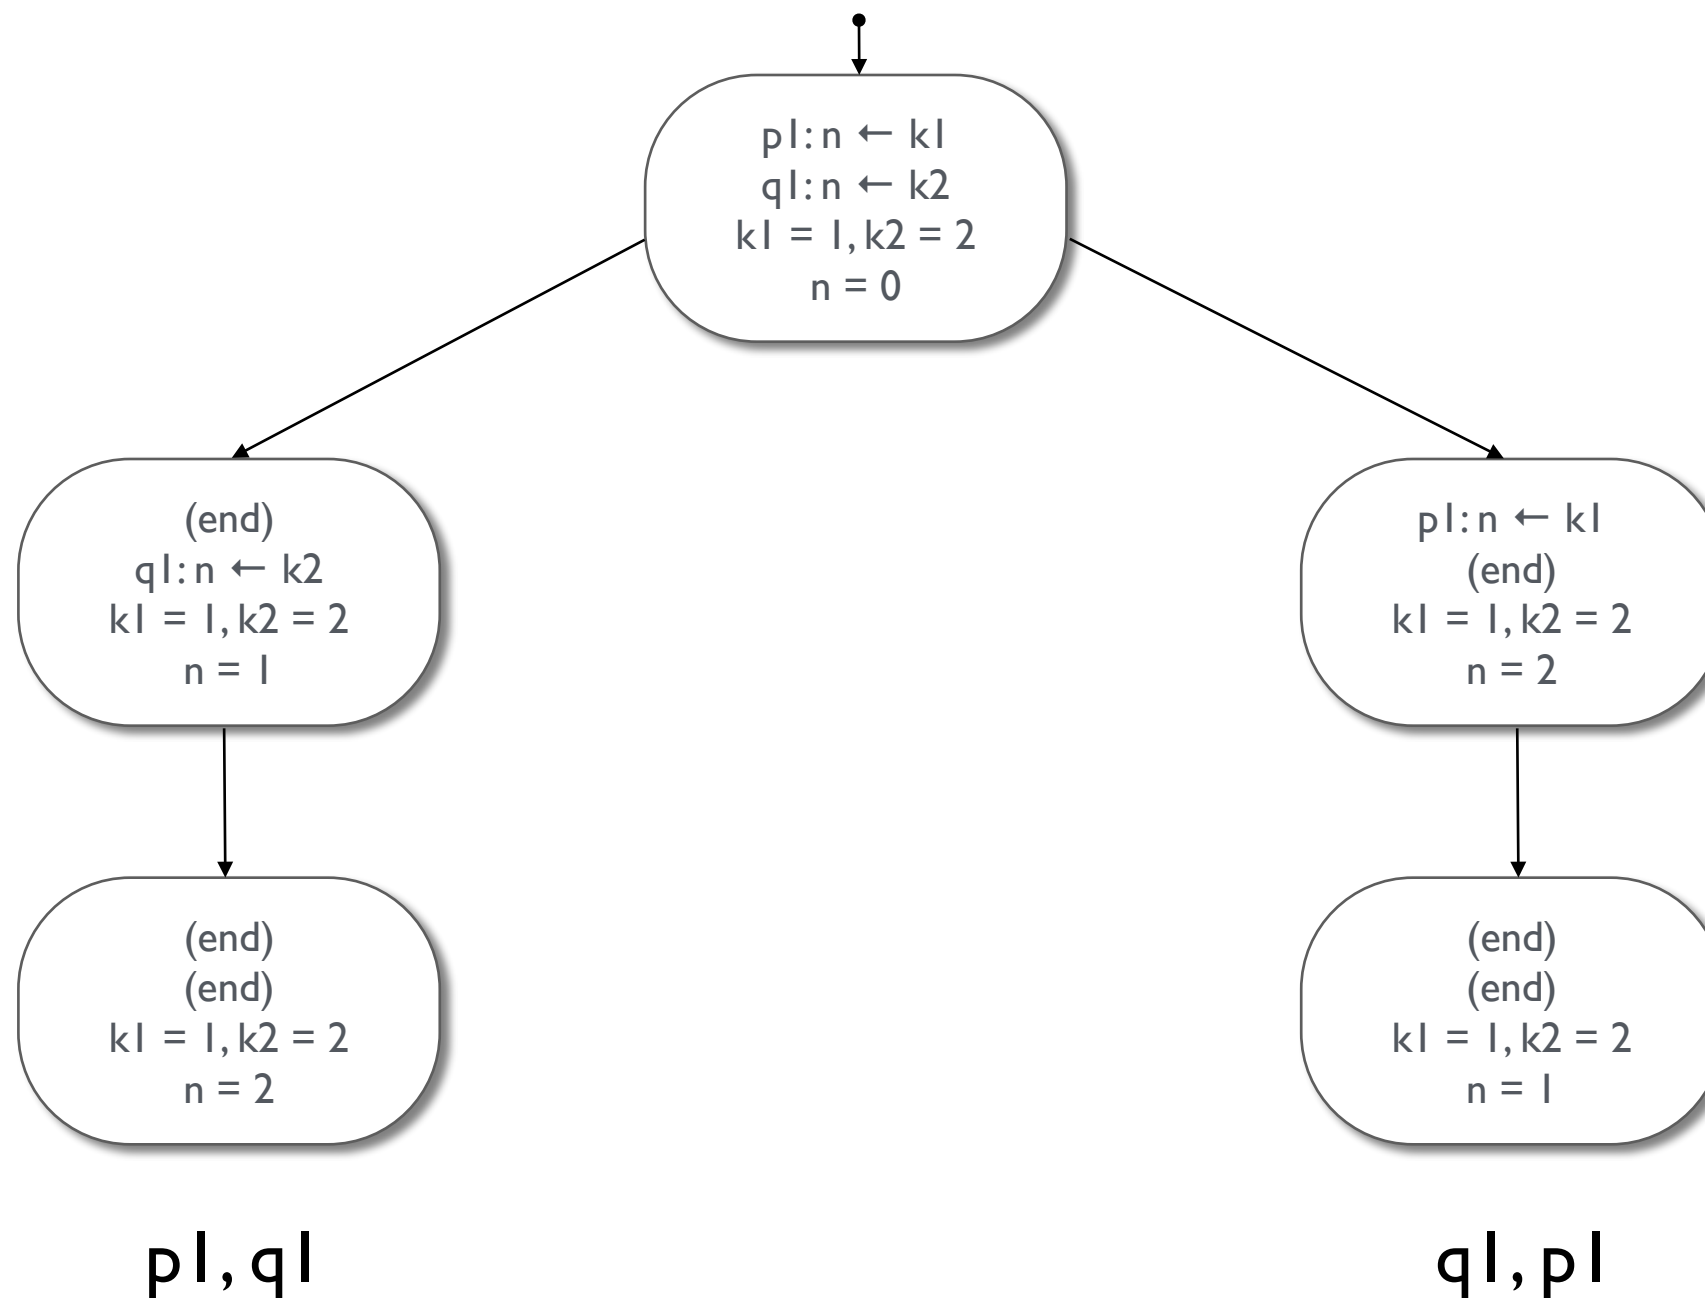
# Simple Concurrent Program (I)

Trivial concurrent program: processes **p** and **q** each do one action that updates global  $n$  with the value of their local variable.

Trivial Concurrent Program	
integer $n \leftarrow 0$ (Globals)	
<b>p</b>	<b>q</b>
integer $k1 \leftarrow 1$ (Locals)	integer $k2 \leftarrow 2$
p1: $n \leftarrow k1$ (Atomic Statements)	q1: $n \leftarrow k2$



# Simple Concurrent Program (2)




# State Diagrams and Scenarios

- We could describe all possible ways a program can execute with a state diagram.
  - There is a transition between  $s_1$  and  $s_2$  (" $s_1:s_2$ ") if executing a statement in  $s_1$  changes the state to  $s_2$ .
  - A state diagram is generated inductively from the starting state.
    - If  $\exists s_1$  and a transition  $s_1:s_2$ , then  $\exists s_2$  and a directed arc from  $s_1:s_2$
- Two states are identical if they have the same variable values and the same directed arcs leaving them.
- A *scenario* is one path through the state diagram.



# Example — Jumping Frogs




- A frog can move to an adjacent stone if it's vacant.
- A frog can hop over an adjacent stone to the next one if that one is vacant.
- No other moves are possible.







# Can we interchange the positions of the frogs?




to




If the frogs can only move “forwards”, can we:



move from above to below?



# Solution Graph



- So, we have a *finite state-transition diagram* of a *finite state machine (FSM)* as a *complete description* of the behaviour of the four frogs, operating concurrently, no matter what they do according to the rules.
- By examining the FSM, we can state properties as definitely holding, i.e. we can prove properties of the system being modelled.

# Solution Graph

- The solution graph makes it clear that this concurrent system—of four frogs that share certain resources—can experience *deadlock*.
- *Deadlock* occurs when the system arrives in a state from which it can not make any transitions (and which is not a desired end-state.)
- *Livelock* (not possible in this system) is when the system can traverse a sequence of states indefinitely without making progress towards a desired end state.



# Proof

- We can prove interesting properties by trying to construct a state diagram describing
  - each possible state of the system and
  - each possible move from one state to another
- We might use a state diagram to show that
  - A property always holds
  - A property never holds in any reachable state
  - A property sometimes holds
  - A property always holds eventually



# State Diagrams for Processes

- A *state* is defined by a tuple of
  - for each process, the label of the statement available for execution.
  - for each variable, its value.
- Q: What is the maximum number of possible states in such a state diagram?



# Abstraction of Concurrent Programming

- A *concurrent program* is a finite set of [sequential] *processes*.
- A process is written using a *finite set of atomic statements*.
- Concurrent program execution is modelled as proceeding by executing a sequence of the atomic statements obtained by *arbitrarily interleaving* the atomic statements of the processes.
- A *computation* [a.k.a. a *scenario*] is one particular execution sequence.





# Atomicity

- We assume that if two operations  $s_1$  and  $s_2$  really happen at the same time, it's the same as if the two operations happened in either order.
- E.g. simultaneous writes to the same memory locations:

Sample	
integer $g \leftarrow 0$ ;	
$p$	$q$
$p! : g \leftarrow 2$ ;	$q! : g \leftarrow 1$

- We assume that the result will be that  $g$  will be 2 or 1 after this program, not, for example, 3.



# Interleaving

- We model a scenario as an arbitrary interleaving of atomic statements, which is somewhat unrealistic.
- For a *concurrent* system, that's OK, it happens anyway.
- For a *parallel* shared memory system, it's OK so long as the previous idea of atomicity holds at the lowest level.
- For a *distributed* system, it's OK if you look at it from an individual node's POV, because either it is executing one of its own statements, or it is sending or receiving a message.
  - Thus *any* interleaving can be used, so long as a message is sent before it is received.



# Level of Atomicity

- The level of atomicity can affect the correctness of a program.

Example: Atomic Assignment Statements	
integer $n \leftarrow 0$ ;	
p	q
p1: $n \leftarrow n+1$ ;	q1: $n \leftarrow n+1$ ;

Value **before** atomic statement on same row

process p	process q	n
<b>p1: <math>n \leftarrow n+1</math>;</b>	q1: $n \leftarrow n+1$ ;	0
(end)	<b>q1: <math>n \leftarrow n+1</math>;</b>	1
(end)	(end)	2

process p	process q	n
p1: $n \leftarrow n+1$ ;	<b>q1: <math>n \leftarrow n+1</math>;</b>	0
<b>p1: <math>n \leftarrow n+1</math>;</b>	(end)	1
(end)	(end)	2



# Different Level of Atomicity

Example: Assignment Statements with one Global Reference	
integer n $\leftarrow$ 0;	
p	q
integer temp	integer temp
p1: temp $\leftarrow$ n	q1: temp $\leftarrow$ n
p2: n $\leftarrow$ temp + 1	q2: n $\leftarrow$ temp + 1



# Alternative Scenarios

process p	process q	n	p.temp	q.temp
<b>p1: temp ← n</b>	q1: temp ← n	0	?	?
<b>p2: n ← temp+1</b>	q1: temp ← n	0	0	?
(end)	<b>q1: temp ← n</b>	1		?
(end)	<b>q2: n ← temp+1</b>	1		1
(end)	(end)	2		

process p	process q	n	p.temp	q.temp
<b>p1: temp ← n</b>	q1: temp ← n	0	?	?
p2: n ← temp+1	<b>q1: temp ← n</b>	0	0	?
<b>p2: n ← temp+1</b>	q2: n ← temp+1	0	0	0
(end)	<b>q2: n ← temp+1</b>	1		0
(end)	(end)	1		



# Atomicity & Correctness

- Thus, the level of atomicity specified affects the correctness of a program
  - We will typically assume that:
    - assignment statements and
    - condition statement evaluations
- are atomic
  - Note: this is not true for programs written in C using concurrency libraries like **pthread**s or similar.



# Concurrent Counting Algorithm

Example: Concurrent Counting Algorithm	
integer $n \leftarrow 0$ ;	
<b>p</b>	<b>q</b>
integer temp	integer temp
p1: do 10 times	q1: do 10 times
p2:    temp $\leftarrow n$	q2:    temp $\leftarrow n$
p3: $n \leftarrow \text{temp} + 1$	q3: $n \leftarrow \text{temp} + 1$

- Increments a global variable  $n$  20 times, thus  $n$  should be 20 after execution.
- But, the program is faulty.
  - Proof: construct a scenario where  *$n$  is 2* afterwards.
- Wouldn't it be nice to get a program to do this analysis?

