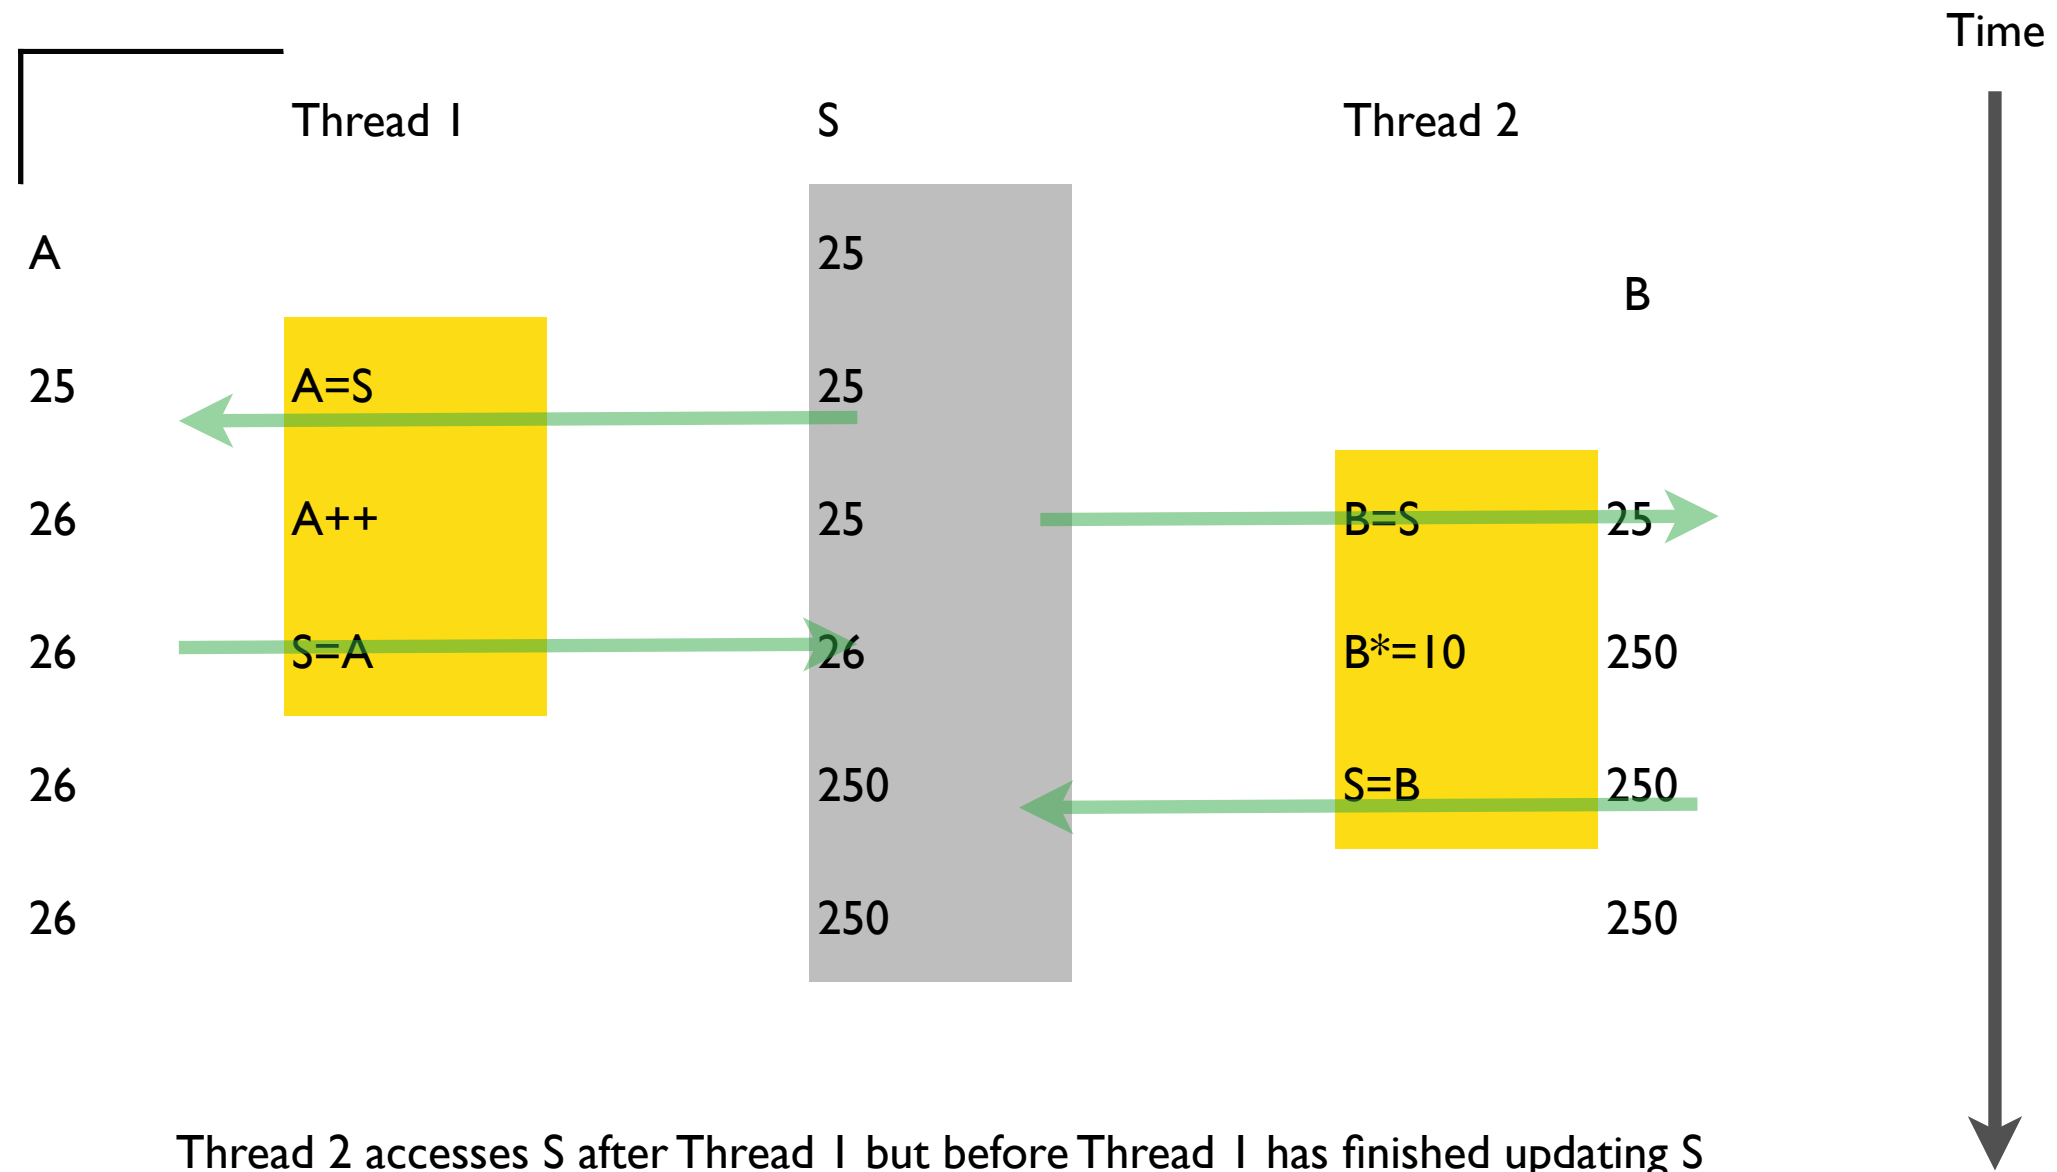


# Synchronisation

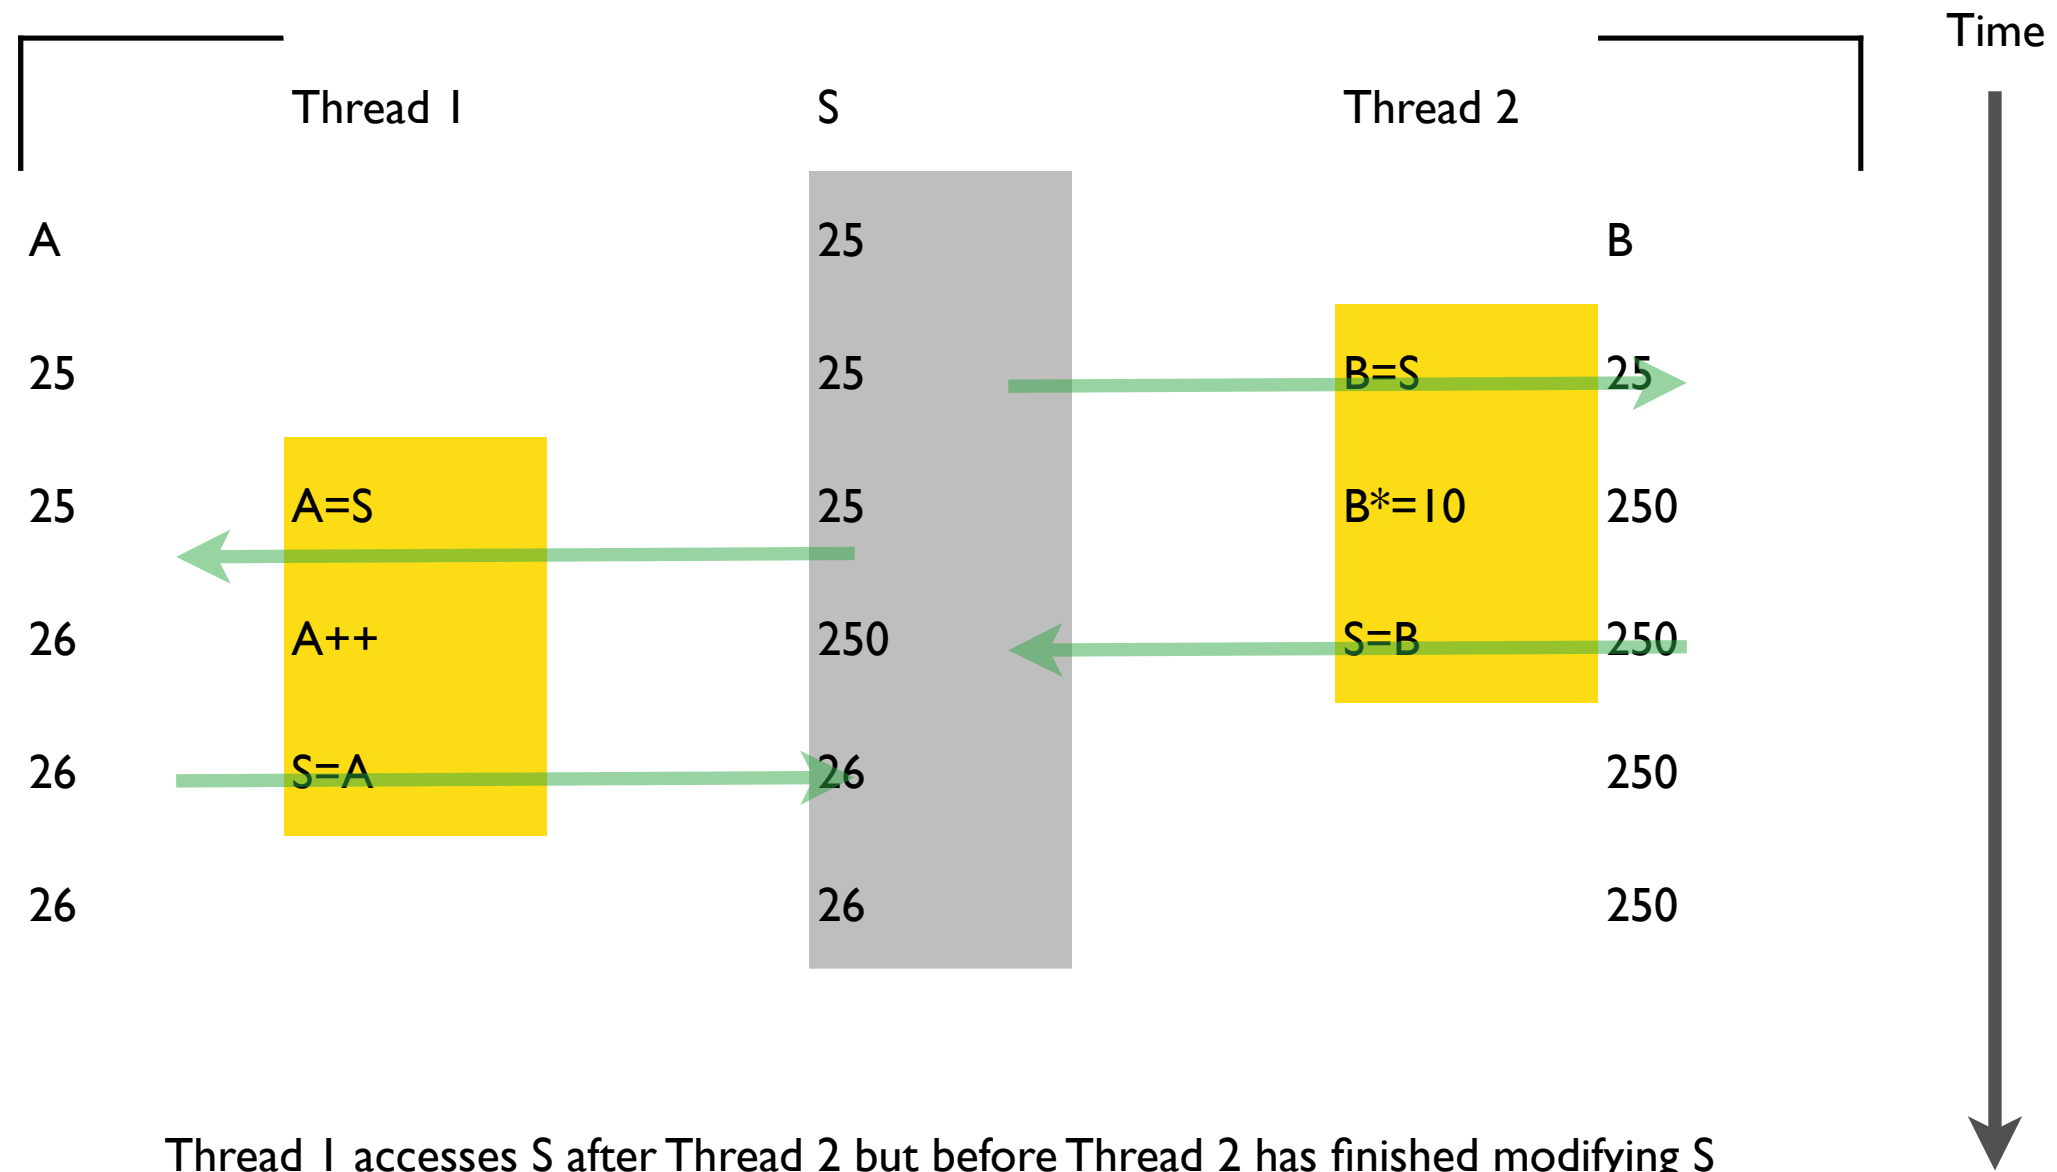
- With every thread we can associate a ‘write-set’:
  - the set of memory locations to which that thread writes.
- We’ve looked at situations where the threads can operate independently -- the ‘write sets’ of the threads don’t intersect.
- Where the write sets intersect, we must ensure that independent thread writes do not damage the data.



# Shared Variable S: Thread 1 *before* Thread 2



# Shared Variable S: Thread 1 *after* Thread 2



# Synchronisation

- Problem: Access to shared resources can be dangerous.
  - These are so-called 'critical' accesses.
- Solution. Critical accesses should be made exclusive. Thus, all critical accesses to a resource are *mutually exclusive*.
- In the example, both threads should have asked for exclusive access before making their updates.
  - Depending on timing, one or the other would get exclusive access first. The other would have to wait to get any kind of access.



# Synchronisation

- A wide variety of algorithms have been devised that ensure synchronisation of competing threads:
  - “Semaphores”, “Monitors”, ...
- All are very tricky to get right
- A helpful way to get simpler guarantees of correct behaviour is to have hardware support.
  - An “atomic” test-and-set instruction  
(read a memory location, check its value, and update the value if needed)
  - An “atomic” operation is one that either performs **ALL** its steps (without interruption) or does **NONE** of them.  
It never does just **SOME** of them.



# Ensuring Mutual Exclusion.

- Mutual Exclusion is a form of thread synchronisation
  - Used to manage access to shared resources (data, IO hardware, shared code, etc.)
- Key idea:
  - There is a "lock" associated with the resource
  - It is locked by a thread when it is using the resource, and unlocked by that thread when done
  - Any other thread must wait for it to be unlocked before proceeding to use that resource.
- accomplished using special "*mutex variables*".

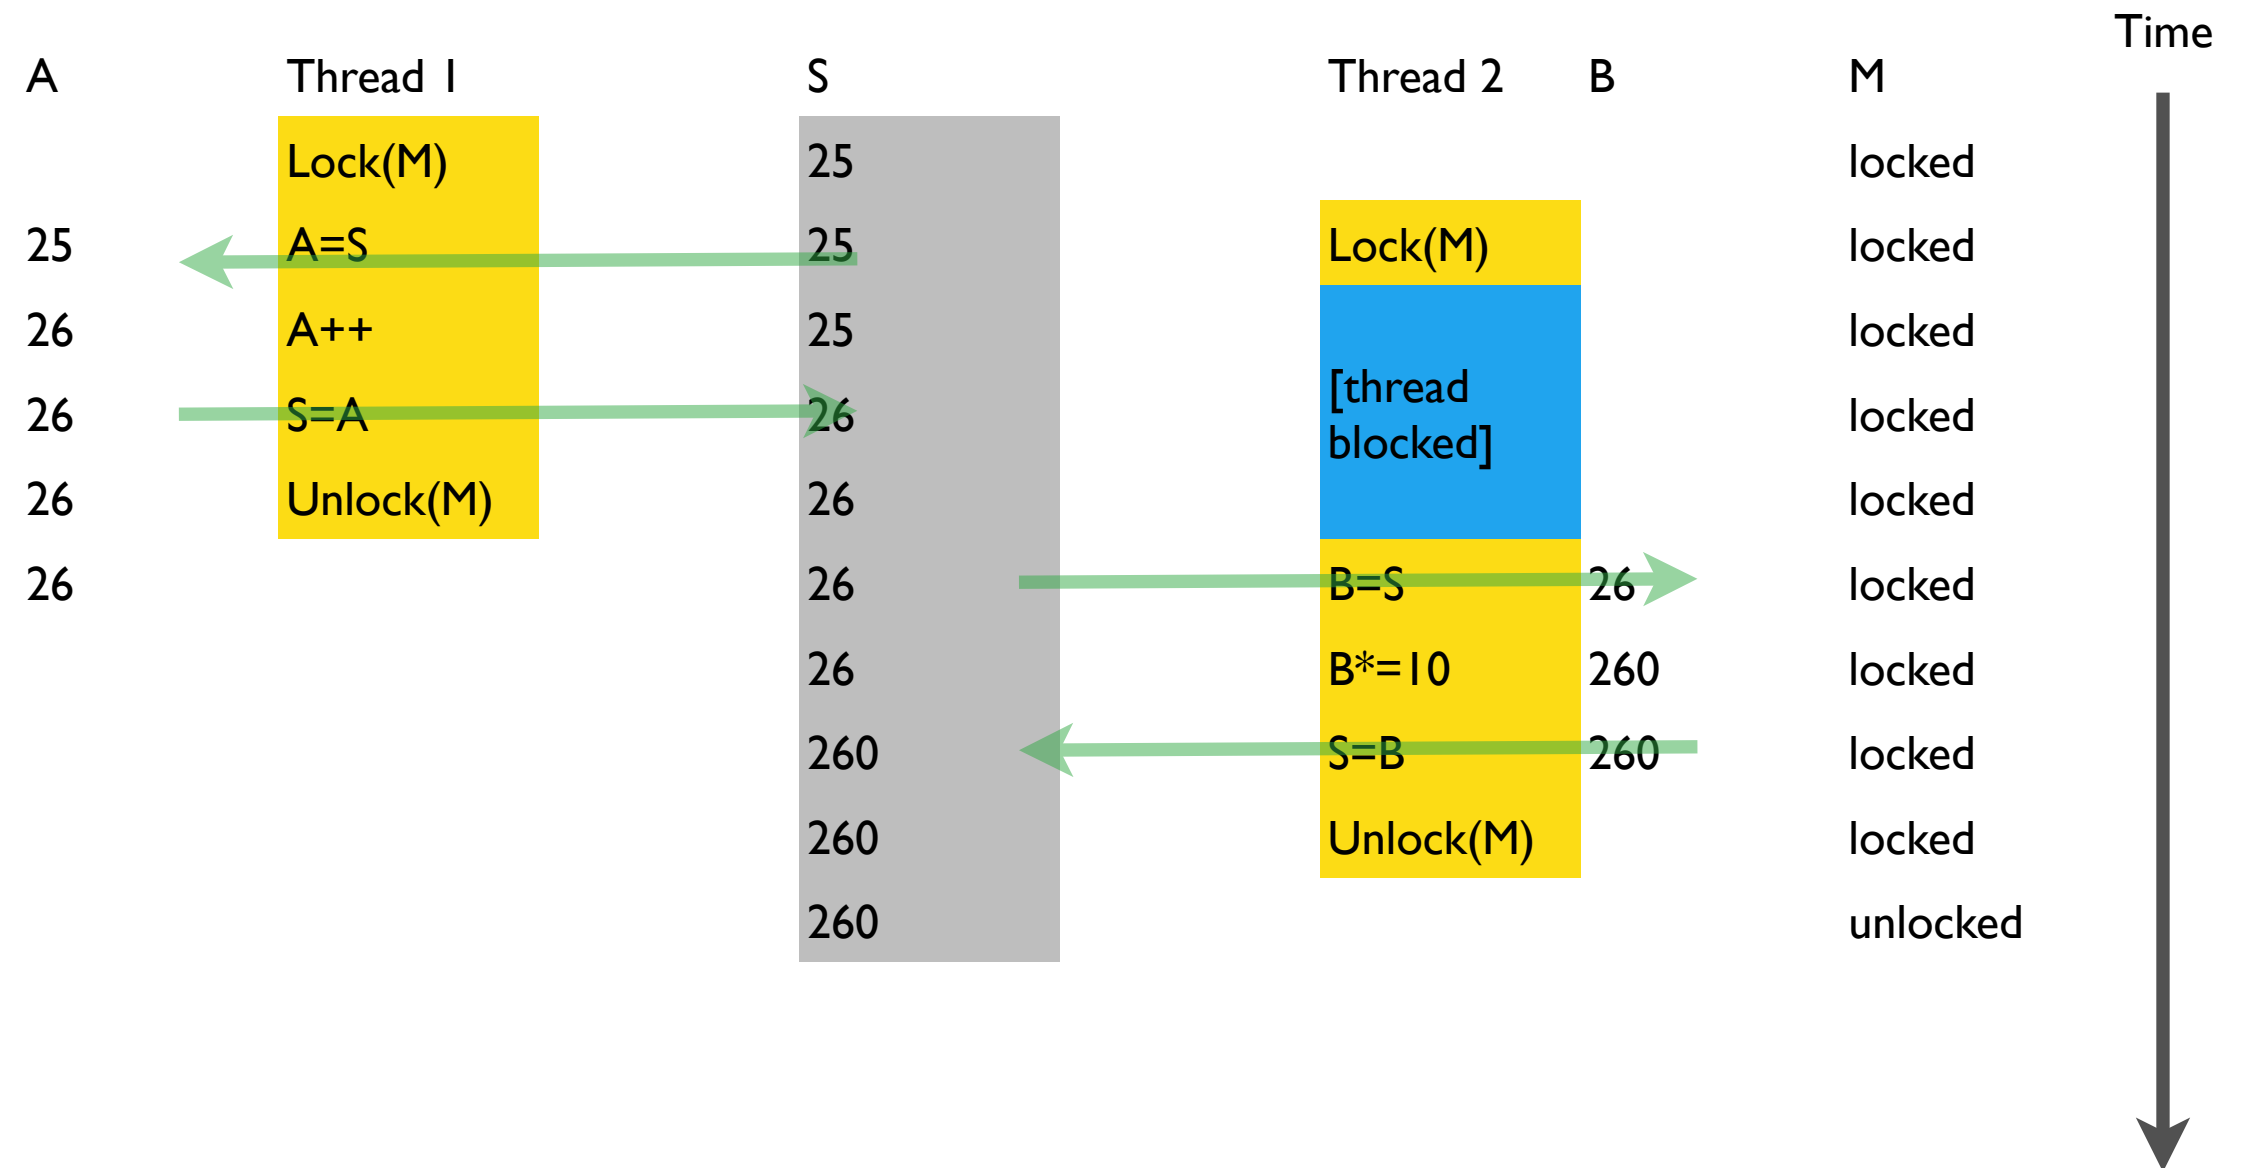


# Accessing a protected resource

- To access a mutex-protected resource, an agent must acquire a lock on the mutex variable.
  - If the mutex variable is unlocked, it is immediately locked and the agent has acquired it. When finished, the agent must unlock it.
  - If the mutex variable is already locked, the agent has failed to acquire the lock -- the protected resource is in exclusive use by someone else.
    - The agent is usually blocked until lock is acquired.
    - A non-blocking version of lock acquisition is available.
- A mutex variable is a complex data object
  - It needs to record if it is locked or unlocked (obviously).
  - It needs to record which thread has acquired the lock.
  - It also needs to keep track of which threads are waiting for it to be unlocked.



# Shared Variable S Protected by Mutex M





# Create pthread mutex variable

- Static:

- `pthread_mutex_t m = PTHREAD_MUTEX_INITIALIZER;`

- Initially unlocked

- Dynamic

- `pthread_mutex_init(<ref to mutex variable>, attributes)`



# Lock and Unlock Mutex

- `pthread_mutex_lock(<mutex variable reference>);`
  - acquire lock or block while waiting
- `pthread_mutex_trylock(<mutex variable reference>);`
  - non-blocking; check returned code
- `pthread_mutex_unlock(<mutex variable reference>);`
  - Should only ever be called by the thread that has the lock!



# Problems

- Voluntary

- Mutexes ‘protect’ code, and only if all thread code uses them properly
- The use of a mutex lock/unlock is a protocol that threads should follow so that code behaves well.
  - It cannot enforce good behaviour
- Other programmers don’t have to use them to get access to the protected resource
  - This is part of the tradeoff. Use processes rather than threads if you want better protection.

- Unfair

- If multiple threads are blocked on a mutex, the order in which they waken up is not guaranteed to be any particular order.



$$\int_0^4 (16 - x^2)dx = 42.6666\dots$$

- We break the interval  $[0..4]$  into a thousand vertical slices ( $H=0.004$ ),
  - use the “trapezoidal rule” to calculate the area of each slice
  - sum all of these areas into an answer variable.
- Sequential solution in `seq-integrate.c` gives an answer of 42.666656
- Concurrent solution in `nom-integrate.c` usually gives the same answer
  - but occasionally gives results like 42.539102 or 42.602656 or 42.602853 or ....
- The problem arise when two threads try to perform  
    `answer = answer + area ;`  
at almost exactly the same time.



# Integration with mutex

- We finally decide to use a mutex (mutex-integrate.c)
- We get the same results as the sequential code.



# mutex-integrate (mutex & thread)

```
pthread_mutex_t  mutex = PTHREAD_MUTEX_INITIALIZER;
```

```
void *IntegratePart(void *i) {
```

```
    double a,b,area;
```

```
    int rc;
```

```
    a = (int)i * H ;
```

```
    b = a + H;
```

```
    area = trapezoid(a,b);
```

```
    // critical section with mutex
```

```
    rc = pthread_mutex_lock(&mutex);
```

```
    checkResults("pthread_mutex_lock()\n", rc);
```

```
    answer=answer+area;
```

```
    rc = pthread_mutex_unlock(&mutex);
```

```
    checkResults("pthread_mutex_lock()\n",rc);
```

```
    pthread_exit(NULL);
```

```
}
```

Body is same as for nom-integrate



# integration example wrap-up

- We can parallelise integration, but mutexes are required
- If we omit them, then errors may occur
- The worst case is when such errors are rare
  - such errors may not be revealed by testing
  - e.g. Mars Rover flash memory bug.
- What about speed?
  - Mutex usage statistics: <http://0pointer.de/blog/projects/mutrace.html>

