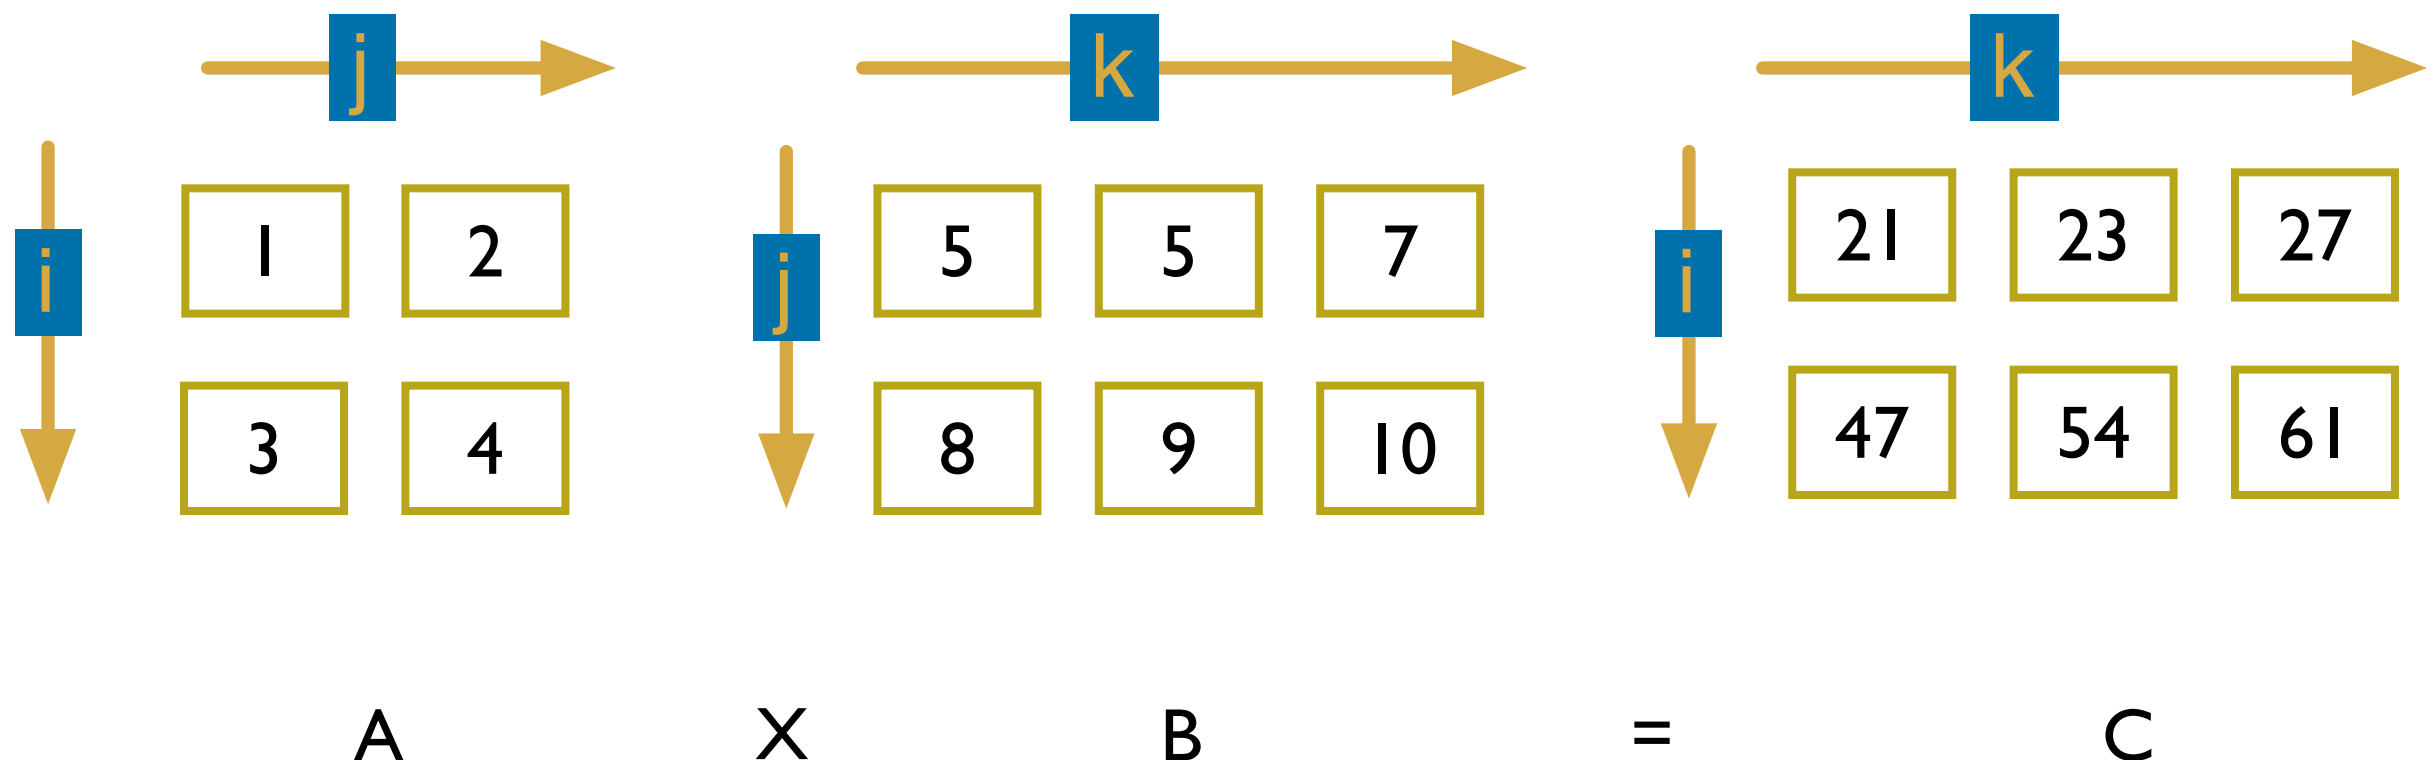


# Removing Unneeded Dependencies in Loops

- Example. We know, intuitively, we can parallelise this a great deal:

- each element in C is independent



# Sample Serial Solution

```
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```



# How can we transform it?

- We can work out how to transform it by locating the dependencies in it.
- Some dependencies are intrinsic to the solution, but
- Some are *artefacts* of the way we are solving the problem;
  - If we can identify them, perhaps we can modify or remove them.



# Try three execution agents:

```
with k = 1;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```

```
with k = 2;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```

```
with k = 3;  
for i = 1 to 2 {  
for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
        sum = sum + a[i,j] * b[j,k];  
    c[i,k] = sum;  
    }  
}
```



# Issues:

- The variable ***sum***, as written, is common to all three programs.
- Solution:
  - Make ***sum*** private to each program to avoid this dependency.



# Try Six Execution Agents

```
with k = 1, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 2, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 3, i=1;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 1, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 2, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```

```
with k = 3, i=2;  
for i = 1 to 2 {  
  for k = 1 to 3 {  
    sum = 0.0;  
    for j = 1 to 2  
      sum = sum + a[1,j] * b[j,k];  
    c[i,k] = sum;  
  }  
}
```



# Summarising:

- We could parallelise the original algorithm with some care:
  - Private Variables to avoid unnecessary dependencies
  - But we may need to combine private results at the end to get a global answer.
- Actually, we could break this ‘Embarrassingly Parallel’ problem into tiny separate pieces; maybe too small. (*How will we know?*)



# How fast can we go?

- We have a sequential program that runs too slow
- We have extra hardware resources that could be used to speed it up.
- How fast can we go?





# Do the math ....

$T(n)$	Time to run program with $n$ parallel processors
$T$	Shorthand for $T(1)$
$S(n)$	Speedup with $n$ processors
$E(n)$	Efficiency of Speedup
$p$	Proportion of $T$ spent executing parallelisable part.
$s$	Speedup possible for parallelisable part



# Speedup and Efficiency

Maximum speedup:

$$T(n) \geq T(1)/n$$

Speedup:

$$S(n) = T(1)/T(n)$$

Efficiency:

$$E(n) = S(n)/n$$



# Program Time and Effective Speedup

- Time to run program without parallelism:

$$T = (1 - p)T + pT$$

- Parallel (effective) speedup vs. processor count:

$$s \leq n$$



# Speedup related to $n$ and $s$ .

Time to run program with speedup  $s$  of parallelisable part:

$$T(n) = (1 - p)T + pT/s = (1 - p + p/s)T$$

Speedup when running program with parallelisable speedup  $s$ :

$$\begin{aligned} S(n) &= T/(1 - p + p/s)T \\ &= 1/(1 - p + p/s) \end{aligned}$$



# Amdahl's Law

$$S(n) = \frac{1}{1 - p + (p/s)} \quad s \leq n$$

$p$  - proportion of single processor time we can parallelise


$n$  - number of parallel processors

$s$  - parallel speedup ratio

$S(n)$  - overall program speedup



# Graphically, Bad News



# Amdahl's Law as $n$ (and $s$ ) get large

$$S(n) = \frac{1}{1 - p + (p/s)} \quad s \leq n$$

Limit as  $n, s$  get very large, so  $p/s \rightarrow 0$ :

$$\frac{1}{1 - p}$$

$$p=0.925, S(n) \longrightarrow 13.333\dots;$$

$$p=0.75, S(n) \longrightarrow 4$$



# Implications

- Even a small fraction of sequential code in a program can seriously interfere with speedup.
  - Note that the code protected by a mutex can only run sequentially!
  - If code has wait a while for a mutex, then that waiting time, has to be added in.
- To maximise performance, inherently sequential code has to be minimised.

