

# How SPIN handles Linear Temporal Logic

- For every LTL predicate, there exists an NDFA that accepts any event sequence that satisfies that predicate.
- Because SPIN is looking for failures, it needs an NDFA that accepts any event sequence that **does not** satisfy that predicate.
- So we generate the NDFA for the **negation** of the predicate, which is then represented as a process that can be written in Promela.
  - This process is called a “never claim”
- This NDFA is then run in parallel with the rest of the Promela model.
  - If it ever enters an “accepting” state, then it has found a violation of the predicate
  - It also signals a failure if an **assert ( )** statement fails.



# Writing LTL in Promela

- According to the books on the reading list:
  - The syntax of LTL in Promela only has variables, and the logical operators.
    - It does not have expressions (boolean or otherwise), such as  $x < 1$  or  $y = z + 42$ .
  - Instead we need to use `#define` to link a variable to such expressions.
  - So `[] ( (x==42) -> <>(y==99) )`

has to be written as something like this:

```
#define answer (x==42)
#define balloons (y==99)
[] (answer -> <>balloons)
```

- However since version 6 of Spin:
  - We can simply write `[] ( (x==42) -> <>(y==99) )`



# From LTL to never claims

- Early versions of Promela/SPIN required the modeller to work out never claims by hand
- Then a version came along that could translate an LTL predicate into a never claim.

```
Promela> spin -f 'p'
never {      /* p */
accept_init:
T0_init:
    do
        :: atomic { ((p)) -> assert(!((p))) }
    od;
accept_all:
    skip
}
```

The output of `spin -f '...'` can be piped into a file which is then included in the Promela model file.

```
Promela> spin -f '<>p'
never {      /* <>p */
T0_init:
    do
        :: atomic { ((p)) -> assert(!((p))) }
        :: (1) -> goto T0_init
    od;
accept_all:
    skip
}
```

```
Promela> spin -f '[ ]p'
never {      /* [ ]p */
accept_init:
T0_init:
    do
        :: ((p)) -> goto T0_init
    od;
}
```

- Current versions allow LTL predicates to be written directly in the Promela file



# Example: a bad mutex solution

- Example of an attempt to solve mutual exclusion
  - It ensures only  $P()$  or  $Q()$  in the critical section at one time
  - Unfortunately, it can deadlock!
    - Easily demonstrated by `spin -run`.
  - How can we use LTL to verify the exclusion property?

Principles of the Spin Model Checker,  
M. Ben-Ari, Sec 4.3, p52

```
1  /* Copyright 2007 by Moti Ben-Ari
2     under the GNU GPL; see readme.txt */
3
4  bool wantP = false, wantQ = false;
5
6  active proctype P() {
7      do :: wantP = true;
8          !wantQ;
9          // critical section
10         wantP = false
11     od
12 }
13
14 active proctype Q() {
15     do :: wantQ = true;
16         !wantP;
17         // critical section
18         wantQ = false
19     od
20 }
```



# mutual exclusion using LTL

- We add a counter **critical**
- Each process:
  - increments it on entry to the critical region
  - decrements it on exit from the critical region
- We define a LTL predicate **msafe** that asserts that **critical** is always less than 2 in any state.
  - If we use **spin -a** we can see the never claim generated, in **\_spin\_nvr.tmp**

Principles of the Spin Model Checker,  
M. Ben-Ari, Sec 5.3.2, pp75-78

```
1  /* Copyright 2007 by Moti Ben-Ari
2     under the GNU GPL; see readme.txt */
3  bool wantP = false, wantQ = false;
4  byte critical;
5  active proctype P() {
6      do :: wantP = true;
7          !wantQ;
8          critical++;
9          // in critical section
10         critical--;
11         wantP = false
12     od
13 }
14 active proctype Q() {
15     do :: wantQ = true;
16         !wantP;
17         critical++;
18         // in critical section
19         critical--;
20         wantQ = false
21     od
22 }
23 ltl msafe { [](critical <=1) }
```



# Other ways ... ?

- We have seen shared-variable concurrency
  - Threads have **non-atomic** read-write access to shared global variables in other threads that are part of the **same** Process.
  - The pthread library provides a discipline for avoiding unwanted interference
  - The Promela modelling we have done has dealt with this approach
- What if we didn't share variables?
  - How can we do this?
  - Would this make all our problems go away?



# Other Ways.

- The main alternative is some form of **message passing**
- Processes/Threads communicate by sending and receiving messages
  - All variables are kept local
- Message passing is the main mechanism used by operating systems to allow communication between Processes
  - Remember, different Processes do not (usually) share any memory.
- Message Passing is the only way to go when doing any networking code
  - Different computers, far apart, cannot share main memory!



# To Synchronise or not to Synchronise?

- There are two forms of message passing: synchronous, and asynchronous
- Synchronous:
  - both sender and receiver wait until the communication has completed
  - If one thread does a send/receive, but another never does a receive/send, then that thread will deadlock!
  - Also known as “Rendezvous Communication”
- Asynchronous:
  - the sender returns immediately, while the message is buffered somewhere
  - the receiver will wait for a message if necessary
  - again the only way to efficiently perform networking
- Hybrid: it is possible to mix the two





# Message Passing in Promela

- Distributed (network) systems consist of nodes connected by communication channels
  - Internet: computers connected by wires, optical fibres, wireless, satellite, and networking hardware running network and communication protocols
- Promela can model this
  - Nodes/Computers are modelled by Promela Processes
  - Network Communication is modelled using Promela Channels



# Channels in Promela

- Declaring a channel: `chan ch = [capacity] of { typename, ..., typename }`
  - capacity is size of (hidden) buffer - can be zero, or positive
  - typename cannot be an array type, but can be a struct that contains an array
  - typename can be `chan` - so we can send/receive communication channels !
- Writing/Sending to a channel: `ch ! expr, ..., expr`
  - the number and type of `expr` must match the types in the declaration for `ch`
- Reading/Receiving from a channel: `ch ? var, ..., var`
  - the number and type of `var` must match the types in the declaration for `ch`



# Simple example: Server + 2\*Client

PSMC, Listing 7.1, p107

- Channel `request` has no buffer
  - Synchronous communication
  - Only time in Promela that two threads perform a step at the exact same time (one does a send-step, the other does a receive step).
- Server loops forever
  - waiting for input on channel request
- Clients make one request and stop.
  - both write to the same channel

```
1  /* Copyright 2007 by Moti Ben-Ari
2     under the GNU GPL; see readme.txt */
3
4  chan request = [0] of { byte };
5
6  active proctype Server() {
7      byte client;
8  end:
9      do
10         :: request ? client
11            -> printf("Client %d\n", client);
12      od
13  }
14
15  active proctype Client0() {
16      request ! 0;
17  }
18
19  active proctype Client1() {
20      request ! 1;
21  }
```



# Buffered Channels

- If channel capacity is non-zero, then we can have asynchronous communication
  - Send puts stuff in buffer, if not full
  - Receive takes stuff, if not empty
- There are builtin predicates to check if a channel buffer is full, empty:
  - `full`, `empty`, `nfull`, `nempty`
  - e.g. `empty(mychan)` is executable/true if the buffer for mychan has no messages,



```

1  /* Copyright 2007 by Moti Ben-Ari under the GNU GPL; see readme.txt */
2
3  chan request = [2] of { byte, chan};
4  chan reply[3] = [2] of { byte };
5
6  bool waiting = false;
7  /* Verify []!waiting to show clients waiting */
8
9  active [2] proctype Server() {
10     byte client;
11     chan replyChannel;
12     do
13         :: empty(request) ->
14             printf("No requests for server %d\n", _pid)
15         :: request ? client, replyChannel ->
16             printf("Client %d processed by server %d\n", client, _pid);
17             replyChannel ! _pid
18     od
19 }
20
21 active [3] proctype Client() {
22     byte server;
23     do
24         :: full(request) ->
25             waiting = true;
26             printf("Client %d waiting for non-full channel\n", _pid)
27         :: request ! _pid, reply[_pid-2] ->
28             reply[_pid-2] ? server;
29             printf("Reply received from server %d by client %d\n", server, _pid)
30     od
31 }

```



# Shared Variable vs. Message Passing

- See <https://wiki.c2.com/?MessagePassingConcurrency>
- See <https://wiki.c2.com/?SharedStateConcurrency>

